



Re-engineering Intelligence Rather Than Starting From Scratch




The Divergence from the Generative Path



The Architectural Contrast: Generative vs. Decision AI

Dimension	Generative LLM	Decision Architecture (Jev)
Core Task	$P(\text{next_token} \mid \text{previous_tokens})$	$P(\text{decision} \mid \text{state, question, possible_decisions})$
Output Format	Unbounded text strings	Strictly bounded probability distributions
Inference Strategy	Sequential sampling (Token 1 $\rightarrow$ 2 $\rightarrow$ 3)	Parallel evaluation of independent variables
Model Interface	~200,000 vocabulary logits	Typed Primitives (Noul, Choice, Score)

Layer 1: Reusing the Semantic Brain



"Even with \$1B, pretraining from zero is an inefficient use of capital. Combine and adapt existing capabilities."

– Diogo Almeida

Layer 2: Shifting the Mathematical Objective

The LLM Way

Based on the evidence,
I believe the claim
should probably...




Changing the
fundamental task
of the network,
not just forcing an
LLM to output
formatted JSON.

The Decision Way

```
{  
  "accept": 0.08,  
  "manual-review": 0.83,  
  "reject": 0.09  
}
```


Layer 3: The Parallel Non-Autoregressive Sampler

Sequential:




Decoding a 200-token explanation.

Parallel:




Evaluating 10 decision questions requires the compute of 1 generation, not 10 sequential generations.


The Data Lab Philosophy: Avoiding the Power Law



The Synthetic Targeting Pipeline




RLCD: Training for Calibrated Confidence



The objective is not merely choosing the correct answer, but ensuring the emitted probability precisely represents the likelihood of correctness.

Replacing the Language Head Interface



The Decision Primitives

NOUL

Operation:
Sigmoid

Output format:
 $P(\text{false})$, $P(\text{true})$

CHOICE

Operation:
Softmax

Output format:
 $P(A)$, $P(B)$, $P(C)$,
 $P(D)$

SCORE

Operation:
Ordered Distribution


Output format:
 $P(\text{level0})$,
 $P(\text{level1})$,
 $P(\text{level2}) \dots$

The Jev Architecture Synthesized

[Pretrained Intelligence]
+ [System-One Architecture]
+ [Parallel Decision Sampling]
+ [Targeted Synthetic Data]
+ [RLCD Calibration] = JEV


A breakthrough achieved through architectural combination
and data systems, not frontier-scale compute.

From Teardown to Blueprint



You do not need a frontier model to utilize this paradigm. You can build a targeted Verifiable Computation Graph (VCG) engine today.

The Blueprint for Your Own Decision Engine



Achieve frontier-level decision architecture
without frontier-level compute.