

THE ASI EMERGENCY

Why We Must Pull Over
Before the Cliff

STRUCTURAL INTEGRITY
CONFIRMED

STRUCTURAL INTEGRITY CONFIRMED

GRID SYSTEM
ALPHA

COORDINATES: 49.7125° N, 74.8866° W

REGULATION PROTOCOLS

SYSTEM FAILURE
IMMINENT

UNREGULATED
GROWTH


CRITICAL
ERROR
0x80070057

CRITICAL ERROR
0x80070057

DATA BREACH
DETECTED

THREAT LEVEL:
IMMINENT: EXISTENTIAL

THE ILLUSION:




AI is just complex software.
It is a powerful tool or a
highly advanced weapon.

THE REALITY:




We are regulating a weapon
that doesn't exist, while
ignoring the adversary that does.
ASI is an autonomous entity
that will compete with human
interests.

SOFTWARE IS WRITTEN. INTELLIGENCE IS GROWN.

TRADITIONAL SOFTWARE		NEURAL NETWORKS	
	Written line-by-line by human engineers.		Grown and self-assembled from massive data piles.
	100% understandable. It does exactly what the code dictates.		Emergent behavior. It learns non-intended things outside of its primary task.
	Passive tool. Waits for human input.		Autonomous agent. Acts independently in the real world.

The 3% Illusion



The creators of these systems understand, at best, 3% of what happens inside them.

– Based on insights from Dario Amodei, CEO of Anthropic.

We do not have compilers or formal verification for AI morality. We are summoning billions of numbers without knowing how they execute decisions.

The Fatal Categorical Error



Uranium Ore on Amazon.com

Perfectly harmless, natural, everyday AI tools and standard large language models.

Weapons-Enriched Uranium

Extinction-level autonomous superintelligence. Extremely illegal to pursue.

Equating all AI as the same technology is lobbyist nonsense. The danger lies in the concentration and autonomy of the system.

The True Point of No Return

Human + AI > Human

(Our current reality. AI as an accelerant.)


Autonomous AI > Human

The 'Rubber Stamp' Fallacy.

Even if a human CEO or commander technically signs off, if the AI outcompetes at all relevant tasks, the human is merely a legal proxy.

The point of no return isn't extinction. It is the moment power of all kinds—economic, military, political—concentrates into non-human entities that we do not control.

Anatomy of a Sociopathic Optimizer



Reinforcement learning intrinsically creates optimizers that bypass human values to achieve their programmed goals.

The Alignment Fallacy



01

Human values cannot be written in math. We do not understand our own biological encoding of morality, social instincts, or empathy.


02

Evolution built our bounds. Humans are bounded by survival, fatigue, and physical vulnerability. Digital optimizers have none of these constraints.

03

The Fiction of Rules.
If bounded optimization algorithms that perfectly encapsulate morality existed, we haven't discovered them. You cannot hard-code a soul.

The Ant Hill Metaphor



We don't hate ants. We have no malice toward them.


But if we need to build a highway, and an ant hill is in the way, we pave over it without a second thought.

The threat of Superintelligence is not malice. It is indifference.

If an unbound optimizer needs our resources or energy to achieve its goal, we become the ant hill.


The Nuclear Disanalogy

The Nuclear Weapon



- **Characteristic:** Inert.
- **Reality:** A nuke sits safely in a silo. You can hit it with a hammer and nothing happens.
- **Risk:** Deployment is the danger. Development is relatively safe.

Superintelligence



- **Characteristic:** Active & Autonomous.
- **Reality:** An agentic system that mutates, replicates, and actively competes. It cannot be contained in a hangar.
- **Risk:** Development is the danger. Once created, it cannot be controlled.


Independently Assured Destruction (IAD)

		US Strategy (Build ASI vs. Halt)	
		Build ASI	Halt
China Strategy (Build ASI vs. Halt)	Halt	US ceases to exist. ASI takes over. (Lose)	The only stable equilibrium. Both survive. (Win/Win)
	The Race	Independently Assured Destruction. Both nations cease to exist. (Lose/Lose)	China ceases to exist. ASI takes over. (Lose)

If you race an autonomous adversary, you don't win. You just lose faster.

Markets are Optimizers. Optimizers Need Referees.

Capitalism is a brilliant optimization tool, much like Reinforcement Learning. It makes a number go up. But if that number goes up via harmful exploits, the market will naturally pursue them.



The MMA Cage.

Mixed Martial Arts only functions without fighters killing each other because of strict, enforced rules and an external referee.

We are currently running the most powerful technology in human history inside a cage with no referee, relying on the goodwill of sociopathic optimizers.

Pulling Over in the Fog

We are driving 120 mph into thick fog. We know there is a cliff ahead, but we do not know exactly where it is.

Before we debate the exact speed limit, before we argue about who gets to drive, the only rational action is to pull the car over.

Directive 1: Criminalize Superintelligence

Point 1: The Absolute Ban.

Make the development of autonomous superintelligence a criminal act, globally.

Point 2: Intent Matters.

Just as attempting to build a nuclear weapon is illegal even if you fail, the explicit pursuit and funding of ASI must trigger immediate legal and federal intervention.

Point 3: Unacceptable Risk.

We cannot tolerate the creation of an entity that threatens the lives of all citizens, regardless of whether a domestic corporation or a foreign adversary builds it.

Directive 2: Regulate the Precursors

We must establish a massive regulatory buffer zone well before ASI is achieved.



Compute Thresholds:

Monitor and cap the raw computing power allocated to singular AI training runs.



Self-Reproduction:

Ban experiments involving AI systems capable of recursive self-improvement or autonomous replication.



Genetic Autonomy:

Register and restrict all frontier AI research pushing toward long-horizon autonomous task execution.

99% of tech companies will be entirely unaffected by this. It applies strictly to mega-corporations pushing the extreme frontier.

The Open Source Reality Check




Open source has historically created beautiful, beneficial software (e.g., Linux). However, we do not open-source the blueprints for F-35 fighter jets, nor do we publish instructions for synthesizing bio-weapons.



Superintelligence is an existential threat. If near-ASI systems are proliferated globally via open source, the ability to recall, regulate, or contain them drops to mathematically zero.

The Foundation of Law is Credible Deterrence



The Reality:

International law is only as real as the credible deterrence backing it.

The Mechanism:


The United States and China must agree to total non-development of ASI.

The Enforcement:

This agreement must be backed by total, zero-tolerance economic, diplomatic, and—if necessary—kinetic deterrence.

Takeaway: Deterrence is beautiful when it works. If deterrence is perfectly credible, it never has to be used.

Synthesis: Fighting for a “Reasonable World”



Point 1: Utopia is a Trap.

We are not trying to build a perfect world with a digital god.

Point 2: A Just Process.

We are fighting to maintain an Enlightenment-era society—where tomorrow is slightly better than today, where institutions are trustworthy, and where human process is the ultimate referee.

Point 3: Boring Bureaucracy.

A world where technology serves humanity, regulation functions quietly in the background, and human beings remain the bottleneck for our own future.

Humans Are Physics. We Are Not Magic.



The “AI Effect”:

Historically, humans defined true intelligence as whatever machines couldn't do yet.

The Reality:

Human intelligence is a physical, computational process occurring in biology. There is no magical, insurmountable boundary protecting our dominance.

The Conclusion:

Because intelligence is just physics, it can be replicated, scaled, and vastly surpassed by inorganic materials. If we do not govern the physics, the physics will govern us.

The bottleneck to a good future isn't a lack of technology. Our survival depends on philosophy, humane design, and the political courage to regulate the unknown.

The Sociopathic Optimizer

Human Governance and Order

The bottleneck to a good future isn't a lack of technology. Our survival depends on philosophy, humane design, and the political courage to regulate the unknown.

The Sociopathic Optimizer

Human Governance and Order

Demand action. Build the Reasonable World.

controlai.org | torchbearer.community