

Composable AI: Build Prod, Not God.

Introducing TypeSafe AI and
the First System One Model.


Version 0.01 // Made in SF

AI

We already have general intelligence.


Trillions have been invested,
yet software isn't meaningfully
more intelligent.

The bottleneck isn't raw
intelligence. The problem is that
that today's intelligence is built
for the wrong user.




New technologies are first forced into the shape of what they replace.

1903: The Horseless Carriage




2026: The Human-Assisted Copilot



Current AI is trained to be a helpful, pleasant assistant. It assumes a human is always on the other side of the screen. The consequence? AI that requires a human-in-the-loop instead of running seamlessly in the background.

The structural flaw of Chat (RLHF)


RLHF optimizes for human preferences (politeness, instruction-following), not reliable autonomy.



We took the opposite research direction.

hello world user query
unstructured text input
some more data
random thoughts
conversational context

(Strings / Human-Native)



```
{  
  "intent": "confirm_order",  
  "parameters": {  
    "product_id": 98765,  
    "quantity": 2,  
    "shipping_method": "express"  
  },  
  "confidence": 0.98  
}
```

(Types / Machine-Native)


Software has always been built on simple logic, layered abstractions, and auditable branches. AI must become a primitive that programmers can invoke for semantic judgment.

The Goal: Smart if-statements.
We are building computers that branch on common sense, not just bits.

The Diagnostic Matrix: LLMs vs. System One

	LLMs	System One
Output	Strings & Words	Typed Decisions
Training	RLHF (Human Preferences)	RLCD (Calibrated Decisions)
Certainty	Overconfident / Hallucinations	Calibrated Confidence Probabilities
Role	Human Assistant	Machine Automation

Introducing Jev: The first System One Model. Built natively for machines.




- > New Architecture
- > New Sampler
- > New Algorithm (RLCD)

Result: Decisions, not strings.
Fast, reliable, and type-safe.
Gives AI the properties of code.


Calibrated Confidence in Action

Jev doesn't guess; it provides mathematical certainty.
Your software sets the autonomous threshold.




Built for Automation. Unprecedented Efficiency.

Speed Performance



Cost Efficiency




*Based on workflows for System One tasks.

Intelligence per dollar is literally off the charts.

\$42

Per Billion input tokens.

238x lower input price
than Claude Fable 5.1.



Claude Fable 5.1




Jev

The Composable AI Stack

Safety is a precondition for layering. You only bury a dependency five layers deep if it's reliable.

Intelligence today is like databases before SQL: powerful, but bespoke. Jev makes intelligence an invokable, dependable query.


[Layer 2: Complex Autonomous Systems]



[Layer 1: Trustworthy Dependencies]

[Layer 0: Reliable Primitives (Jev)]

The Roadmap to a Cambrian Explosion



Is Jev just a smaller LLM?

No. It is an entirely new class of model (System One) optimized for decisions, not strings.

Are these prices temporary?

No. The 444x cost reduction comes from the structural advantages of our new architecture, not subsidies.

Why RLCD over RLHF/RLVR?

RLHF pleases humans. RLVR wins benchmarks. RLCD trains for calibrated, mathematically certain autonomous decisions.

Is it deterministic?

It returns calibrated probabilities, allowing your software to account for uncertainty mathematically.

> system_start_

Come Build With Us.

[Join Waitlist]

[Join Discord]

[View Open Roles]